

Group 2 Feedback: Precondition Report

MVP scope and feasibility

Your proposed scope is substantial and potentially appropriate for a team of eight. However, it combines several technically demanding capabilities, including URL-based document retrieval, document parsing, source-grounded AI analysis, authentication, saved history, AI chat, and the Data Footprint Map. The concern is not whether the team has enough work, but whether all these components can be integrated, secured, tested, and evaluated within one semester.

In Design Portfolio I, clearly distinguish between:

1. The committed end-to-end MVP that the team must deliver and verify.
2. Secondary capabilities to be added only after the core workflow functions reliably.
3. Stretch goals that are not required for MVP completion.

The AI analysis may remain central because it directly addresses the project's problem. For example, the committed MVP could allow a user to submit a supported public agreement, extract selected privacy categories, receive plain-language findings linked to supporting passages, save the analysis, and view a basic Data Footprint summary.

Document-specific AI chat may also remain in the MVP if the team identifies the additional user need it serves and explains how its responses will be grounded and evaluated. Broader URL or document support beyond the stated machine-readable web and text agreements, richer Data Footprint visualizations, and comparison of saved scans may be treated as secondary capabilities or stretch goals.

Your DP I must show how the committed MVP will provide meaningful, integrated work for all eight members. The requirement of at least six meaningful pull requests per student is a minimum contribution expectation; it does not demonstrate that the entire proposed scope is feasible. The project plan must reserve sufficient time for integration, security, automated testing, AI-output evaluation, and improvement based on feedback.

Data, Privacy, Security, and Safety

The Data, Privacy, Security, and Safety section of your Precondition Report is blank. You do not need to resubmit the Precondition Report, but you must address these matters in DP I through your scope, requirements, architecture, and risk-mitigation decisions.

For example, consider whether the committed MVP should:

- Accept only publicly available, machine-readable agreements and warn users not to submit confidential documents or personal information.
- Send only the document text and document-related question to Claude, excluding names, email addresses, account identifiers, and authentication information.

- Store only the minimum data needed, such as the submitted URL, document metadata, findings, citations, and scan history.
- Allow users to delete saved analyses and associated records.
- Ensure that authenticated users can access only their own saved information.
- Keep AI credentials and other secrets exclusively on the server.
- Restrict document retrieval to safe URLs and block private, local, or internal network addresses.
- Limit document size, request duration, redirects, and API usage.
- Label ambiguous or unsupported AI findings rather than presenting them as facts.
- Define how the system will respond when the AI service is unavailable or when its cost or usage limit is reached.

These are examples, not predetermined requirements. Your team must determine which controls are appropriate for the agreed MVP and express them as specific, testable functional or non-functional requirements in DP I.

Measurable AI-evaluation criterion

The extraction-quality statement, “the system achieves high on a manually labeled evaluation set,” is incomplete because it does not identify an extraction metric or an acceptable threshold.

The following statement about users answering questions at least 15 percentage points more accurately is a separate comparative user-evaluation measure. It does not complete the extraction-quality criterion.

In DP I, define what will be evaluated, the size and composition of the labeled test set, how the reference labels will be produced, the metric to be calculated, and the minimum acceptable result.

For example:

AI extraction quality: On a manually reviewed test set of supported privacy policies, the system shall achieve a macro-averaged F1 score of at least 0.80 when identifying the selected privacy-information categories.

Verification method: Compare the system’s extracted categories with the manually established reference labels and calculate precision, recall, and F1 score for each category and across the complete test set.

This is only an example. Your team must select and justify the final metrics, thresholds, and test-set size based on the committed MVP.

User or stakeholder evaluation

Your semester plan includes user testing during Weeks 9–11, and the Precondition Report instructions require user or stakeholder evaluation. However, Preliminary Success Criterion 4 labels this evaluation as optional. This is inconsistent with both your semester plan and the user study required in DP II.

Resolve this inconsistency in DP I by providing a preliminary plan for the DP II user study. Identify the intended participants or stakeholders, the tasks or features they will evaluate, what the team expects to observe or measure, and how the findings will inform subsequent design or implementation changes. Include the planned study in the project schedule and connect it to an appropriate success criterion and verification method.

At this stage, you do not need to provide the complete study protocol or results. The detailed procedure, execution, findings, and resulting project changes will be reported in DP II.

Retain the proposed 15-percentage-point improvement over reading the original agreement only if the team plans to conduct a suitable comparative user study. Otherwise, replace it with a smaller and more feasible criterion that still provides useful evidence about comprehension, usability, or stakeholder acceptance.

Software verification

The report identifies GitHub Actions and proposes formatting and linting checks, unit tests, integration tests, build validation, and later end-to-end testing. Therefore, the issue is not that testing and CI are absent. The remaining gap is that the preliminary success criteria do not translate this testing approach into measurable verification expectations for the committed MVP.

In DP I, connect the selected testing tools and CI workflow to acceptance criteria for the major functional requirements and to measurable non-functional requirements.

For example, consider specifying that:

- Every pull request must pass the agreed linting, build, unit-test, and integration-test checks before it is merged.
- Automated tests must verify the core workflow from document submission through retrieval, parsing, analysis, and presentation of source-grounded results.
- The system must detect and report unsupported, inaccessible, malformed, or oversized documents without crashing or saving a misleading result.
- Authentication and authorization tests must confirm that users cannot access another user's saved analyses.
- Failure or unavailability of the AI provider must produce a clear error and must not create an incomplete analysis that appears successful.

These are examples. Select verification expectations that correspond directly to the final MVP and state how each one will be checked.

Document completeness and collective review

The report contains several identifiable editing problems:

- The backend description is repeated.
- The three template placeholders under Preliminary Exclusions remain in the document.
- Success Criterion 5 is blank.

- The phrase “too lazy to read EULAs” is inappropriate for a formal technical document.

A more professional description would refer to users who lack the time or expertise to interpret lengthy agreements.

These issues suggest that the assembled report did not receive a sufficiently careful final review. However, the document alone does not establish which individual members contributed or participated.

DP I must be reviewed as one coherent team submission before it is submitted. Assign responsibility for checking completeness, consistency, professional language, duplicated material, and alignment among the MVP, requirements, design, evaluation, and project plan. The entire team remains responsible for the quality of the final submission.

Team coordination and meeting records

Please follow the separate team-coordination feedback and use the new meeting-minutes template shared with the class beginning with your next meeting.

Priorities for Design Portfolio I

In DP I, carry these corrections into the corresponding sections:

1. Clearly distinguish the committed MVP from secondary capabilities and stretch goals in the project scope and requirements.
2. Address data handling, privacy, security, safety, AI-provider risks, and safe URL retrieval in the non-functional requirements, architecture, and risk analysis.
3. Replace incomplete or inconsistent success criteria with justified measures, thresholds, and verification methods.
4. Connect the proposed CI and testing approach to functional requirements, acceptance criteria, and the verification plan.
5. Identify the intended participants or stakeholders for the evaluation.
6. Record individual responsibilities in the project plan and use the new meeting-minutes template to track assignments and progress.

You do not need to resubmit the Precondition Report. Carry these corrections and decisions forward into Design Portfolio I and begin working on it immediately. Before submission, review DP I as one complete document for accuracy, consistency, completeness, and professional language.